

Kaichun Mo

Mailing Address
10307 NE 126TH PL
Kirkland, WA 98034

Email: daerduomkch@gmail.com
Tel: (+1) 650-521-6801
Homepage: <https://kaichun-mo.github.io>
Links: *Google Scholar LinkedIn*

RESEARCH INTERESTS

- Frontier Generative AI Foundation Models (Video-Gen/VLM/3D-Gen);
- Frontier Robotic Foundation Models (VLA/Diffusion Policy);
- 3D Computer Vision, Geometry Learning, and 3D Deep Learning.

EMPLOYMENT

NVIDIA Research, USA
Research Scientist

Aug 2022 – Now

- **Managers:** Ming-Yu Liu;
- **Lab:** Deep Imagination Research.
- **Project Highlights:**
 - Cosmos world foundation model platform for physical AI;
 - Developing the next version of Cosmos-Transfer, Cosmos-Interactive, Cosmos-Reason models.
- **Manager:** Dieter Fox;
- **Lab:** Seattle Robotics Lab.
- **Project Highlights:**
 - Robotic post-training for video generation models;
 - 3D point cloud world model;
 - Agentic VLA models for robot manipulation;
 - Large-scale robot data collection in simulation.

EDUCATION

Stanford University, USA
Ph.D., Computer Science

Sept 2016 – June 2022

- **Advisor:** Prof. Leonidas Guibas;
- **Lab:** Stanford Geometry Lab and Stanford Artificial Intelligence Lab;
- **Disseration:** Learning Compositional and Actionable Visual Representations for 3D Shapes.
- **Project Highlights:**
 - PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation;
 - PartNet: A Large-scale Benchmark for Fine-grained and Hierarchical Part-level 3D Object Understanding;
 - StructureNet: Hierarchical Graph Networks for 3D Shape Generation;
 - SAPIEN: A SimulAteD Part-based Interactive ENvironment;
 - Where2Act: From Pixels to Actions for Articulated 3D Objects.

Shanghai Jiao Tong University, China
B.S.E., Computer Science, ACM Honored Class, Zhiyuan College

Sept 2012 – July 2016

- ACM Honored Class is a highly selective class (top 33 students) in SJTU.
- GPA: 3.96/4.30 (91.87/100) Rank: **1/33**.

PUBLICATIONS

Shengyi Qian, **Kaichun Mo**, Valts Blukis, David Fouhey, Dieter Fox and Ankit Goyal, *3D-MVP: 3D Multiview Pretraining for Robotic Manipulation*, CVPR 2025

Dandan Shan, **Kaichun Mo**, Wei Yang, Yu-Wei Chao, David Fouhey, Dieter Fox and Arsalan Mousavian, *Slot-Level Robotic Placement via Visual Imitation from Single Human Video*, FMEA Workshop @ CVPR 2025.

Yian Wang, Bingjie Tang, Chuang Gan, Dieter Fox, **Kaichun Mo**, Yashraj Narang and Iretoiyo Akinola, *MatchMaker: Automated Robotic Assembly Asset Generation for Policy Learning in Simulation*, ICRA 2025

Zoey Chen, Aaron Walsman, Marius Memmel, **Kaichun Mo**, Alex Fang, Karthikeya Vemuri, Alan Wu, Dieter Fox and Abhishek Gupta, *URDFormer: Constructing Interactive Realistic Scenes from Real Images via Simulation and Generative Modeling*, RSS 2024

Razvan Pasca, Alexey Gavryushin, Muhammad Hamza, Yen-Ling Kuo, **Kaichun Mo**, Luc Van Gool, Otmar Hilliges and Xi Wang, *Summarize the Past to Predict the Future: Natural Language Descriptions of Context Boost Multimodal Object Interaction Anticipation*, CVPR 2024.

Yichen Li, **Kaichun Mo**, Yueqi Duan, He Wang, Jiequan Zhang, Lin Shao, Wojciech Matusik and Leonidas J. Guibas, *Category-Level Multi-Part Multi-Joint 3D Shape Assembly*, CVPR 2024.

Jia-Mu Sun, Jie Yang, **Kaichun Mo**, Yu-Kun Lai, Leonidas J. Guibas and Lin Gao, *HAISOR: Human-Aware Indoor Scene Optimization via Deep Reinforcement Learning*, ACM Transactions on Graphics (SIGGRAPH 2024).

Chuanruo Ning, Ruihai Wu, Haoran Lu, **Kaichun Mo** and Hao Dong, *Where2Explore: Few-shot Affordance Learning for Unseen Novel Categories of Articulated Objects*, NeurIPS 2023.

Yi Li, Muru Zhang, Markus Grotz, **Kaichun Mo** and Dieter Fox, *STOW: Discrete-Frame Segmentation and Tracking of Unseen Objects for Warehouse Picking Robots*, CoRL 2023.

Boxiao Pan, Bokui Shen*, Davis Rempe*, Despoina Paschalidou, **Kaichun Mo**, Yanchao Yang and Leonidas J. Guibas, *COPLOT: Human Collision Prediction and Localization from Multi-view Egocentric Videos*, ICCV 2023

Yijia Weng, **Kaichun Mo**, Ruoxi Shi, Yanchao Yang and Leonidas J. Guibas, *Toward Learning Geometric Eigen-Lengths Crucial for Fitting Tasks*, ICML 2023.

Xiaomeng Xu*, Yanchao Yang*, **Kaichun Mo**, Boxiao Pan, Li Yi and Leonidas J. Guibas, *JacobiNeRF: NeRF Shaping with Mutual Information Gradients*, CVPR 2023.

Yan Zhao*, Ruihai Wu*, Zhehuan Chen, Yourong Zhang, Qingnan Fan, **Kaichun Mo** and Hao Dong, *DualAfford: Learning Collaborative Visual Affordance for Dual-gripper Object Manipulation*, ICLR 2023.

Lin Gao, Jia-Mu Sun, **Kaichun Mo**, Yu-Kun Lai, Leonidas J. Guibas and Jie Yang, *SceneHGN: Hierarchical Graph Networks for 3D Indoor Scene Generation with Fine-Grained Geometry*, TPAMI 2023.

Jeonghyun Kim, **Kaichun Mo**, Minhyuk Sung* and Woontack Woo*, *SegStruct: The Interplay between Part Segmentation and Structure Inference for 3D Shape Parsing*, WACV 2023.

Kaichun Mo*, Yian Wang*, Ruihai Wu*, Jiaqi Ke, Qingnan Fan, Leonidas J. Guibas and Hao Dong, *AdaAfford: Learning to Adapt Manipulation Affordance for 3D Articulated Objects via Few-shot Interactions*, ECCV 2022.

Yang Zheng, Yanchao Yang, **Kaichun Mo**, Jiaman Li, Tao Yu, Yebin Liu, Karen Liu and Leonidas J. Guibas, *GIMO: Gaze-Informed Human Motion Prediction in Context*, ECCV 2022.

Yining Hong, **Kaichun Mo**, Li Yi, Leonidas J. Guibas, Antonio Torralba, Joshua Tenenbaum and Chuang Gan, *Fixing Malfunctional Objects With Learned Physical Simulation and Functional Prediction*, CVPR 2022.

Kaichun Mo*, Jie Yang*, Yu-Kun Lai, Leonidas J. Guibas and Lin Gao, *DSG-Net: Learning Disentangled Structure and Geometry for 3D Shape Generation*, ACM Transactions on Graphics (presented at SIGGRAPH 2022).

Kaichun Mo*, Qi Li*, Yanchao Yang, Hang Zhao and Leonidas J. Guibas, *IFR-Explore: Learning Inter-object Functional Relationships in 3D Indoor Scenes*, ICLR 2022.

Kaichun Mo*, Ruihai Wu*, Yan Zhao*, Zizheng Guo, Yian Wang, Tianhao Wu, Qingnan Fan, Xuelin Chen, Leonidas J. Guibas and Hao Dong, *VAT-Mart: Learning Visual Action Trajectory Proposals for Manipulating 3D ARTiculated Objects*, ICLR 2022. (**awards:** Youth Outstanding Paper Award at WAIC 2025)

ChuanYu Pan*, Yanchao Yang*, **Kaichun Mo**, Yueqi Duan and Leonidas J. Guibas, *Object Pursuit: Building a Space of Objects via Discriminative Weight Generation*, ICLR 2022.

Kaichun Mo, Yuzhe Qin, Fanbo Xiang, Hao Su and Leonidas J. Guibas, *O2O-Afford: Annotation-Free Large-Scale Object-Object Affordance Learning*, CoRL 2021.

Shuo Cheng, **Kaichun Mo** and Lin Shao, *Learning to Regrasp by Learning to Place*, CoRL 2021.

Kaichun Mo, Leonidas J. Guibas, Mustafa Mukadam, Abhinav Gupta and Shubham Tulsiani, *Where2Act: From Pixels to Actions for Articulated 3D Objects*, ICCV 2021.

Jialei Huang*, Guanqi Zhan*, Qingnan Fan, **Kaichun Mo**, Lin Shao, Baoquan Chen, Leonidas J. Guibas and Hao Dong, *Generative 3D Part Assembly via Dynamic Graph Learning*, NeurIPS 2020.

He Wang*, Zetian Jiang*, Li Yi, **Kaichun Mo**, Hao Su and Leonidas J. Guibas, *Rethinking Sampling in 3D Point Cloud Generative Adversarial Networks*, CVPR 2021 Workshop "Learning to generate 3D Shapes and Scenes".

Kaichun Mo*, Yichen Li*, Lin Shao, Minhyuk Sung and Leonidas J. Guibas, *Learning 3D Part Assembly from a Single Image*, ECCV 2020.

Kaichun Mo, He Wang, Xinchun Yan and Leonidas J. Guibas, *PT2PC: Learning to Generate 3D Point Cloud Shapes from Part Tree Conditions*, ECCV 2020.

Kaichun Mo*, Paul Guerrero*, Li Yi, Hao Su, Peter Wonka, Niloy Mitra and Leonidas J. Guibas, *StructEdit: Learning Structural Shape Variations*, CVPR 2020.

Fanbo Xiang, Yuzhe Qin, **Kaichun Mo**, Yikuan Xia, Hao Zhu, Fanchen Liu, Minghua Liu, Hanxiao Jiang, Yifu Yuan, He Wang, Li Yi, Angel Chang, Leonidas Guibas and Hao Su, *SAPIEN: A Simulated Part-based Interactive ENvironment*, CVPR 2020 (Oral).

Tiange Luo, **Kaichun Mo**, Zhiao Huang, Jiarui Xu, Siyu Hu, Liwei Wang, Hao Su, *Learning to Group: A Bottom-Up Framework for 3D Part Discovery in Unseen Categories*, ICLR 2020.

Kaichun Mo*, Paul Guerrero*, Li Yi, Hao Su, Peter Wonka, Niloy Mitra and Leonidas J. Guibas, *Struc-*

tureNet: Hierarchical Graph Networks for 3D Shape Generation, ACM Transactions on Graphics (SIGGRAPH Asia 2019).

Kaichun Mo, Shilin Zhu, Angel X.Chang, Li Yi, Subarna Tripathi, Leonidas J. Guibas and Hao Su, *PartNet: A Large-scale Benchmark for Fine-grained and Hierarchical Part-level 3D Object Understanding*, CVPR 2019. (**featured in:** IEEE Spectrum, The Robot Report, etc.)

Charles R. Qi, Hao Su, **Kaichun Mo**, and Leonidas J. Guibas, *PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation*, CVPR 2017 (Oral). (**awards:** 2025 Frontiers of Science Award)

Yujun Li, **Kaichun Mo** and Haishan Ye, *Accelerating Random Kaczmarz Algorithm Based on Clustering Information*, AAAI 2016.

MANUSCRIPTS

NVIDIA, *Cosmos World Foundation Model Platform for Physical AI*

Alexey Gavryushin, Florian Redhardt, Gaia Di Lorenzo, Luc Van Gool, Marc Pollefeys, **Kaichun Mo** and Xi Wang, *SIGHT: Single-Image Conditioned Generation of Hand Trajectories for Hand-Object Interaction*, arXiv:2503.22869 [cs.CV].

Mingxin Yu*, Lin Shao*, Zhehuan Chen, Tianhao Wu, Qingnan Fan, **Kaichun Mo** and Hao Dong, *RoboAssembly: Learning Generalizable Furniture Assembly Policy in a Novel Multi-robot Contact-rich Simulation Environment*, arXiv:2112.10143 [cs.RO].

Songfang Han, Jiayuan Gu, **Kaichun Mo**, Li Yi, Siyu Hu, Xuejin Chen and Hao Su, *Compositionally Generalizable 3D Structure Prediction*, arXiv:2012.02493 [cs.CV].

Kaichun Mo, Haoxiang Li, Zhe Lin and Joon-Young Lee, *The AdobeIndoorNav Dataset: Towards Deep Reinforcement Learning based Real-world Indoor Robot Visual Navigation*, arXiv:1802.08824 [cs.RO].

Luo Luo, **Kaichun Mo**, Zhihua Zhang, *Blockwise Matrix Completion for Image Colorization*, technical report, 2015.

OTHER RESEARCH EXPERIENCES

Research Intern June 2020 – Sep 2020

Facebook AI Research, Facebook Inc., USA

- Mentors: Shubham Tulsiani, Mustafa Mukadam and Prof. Abhinav Gupta.
- Object-centric segmentation, affordance and actionable information for robotic manipulation.

Visiting Graduate July 2019 – Sep 2019

Su Lab, University of California, San Diego (UCSD), USA

- Advisor: Prof. Hao Su;
- 3D weakly-supervised learning for shape part and structure discovery;
- Object-centric physical simulation for 3D robotic vision, control and planning.

Research Intern June 2018 – Sep 2018

Machine Intelligence group, Autodesk Research, Autodesk Inc., USA

- Mentor: Mike Haley;
- Unsupervised shape structure induction from a collection of 3D shapes.

Research Intern June 2017 – Sep 2017

Imagination Lab, Adobe Research, Adobe System Inc., USA

- Mentors: Haoxiang Li, Joon-Young Lee, Zhe Lin and Ersin Yumer;
- Autonomous robot indoor navigation using Reinforcement Learning.

Rotation Student

April 2017 – June 2017

Computer Graphics Lab, Stanford University, USA

- Advisor: Prof. Doug James;
- Sound synthesis using 3D CAD ShapeNet models.

Rotation Student

Jan 2017 – April, 2017

Computational Vision and Geometry Lab, Stanford University, USA

- Advisor: Prof. Silvio Savarese;
- Multi-agent collaboration and theory of mind using Reinforcement Learning.

Exchange Research Scholar

July 2015 – Dec 2015

Graphics and Vision Lab, Cornell University, USA

- Advisor: Prof. Kavita Bala;
- Algorithmic propagation of material properties from 2D images to 3D ShapeNet models.

Research Assistant

June 2014 – June 2016

Brain-like Computing and Machine Intelligence Lab, Shanghai Jiao Tong University, China

- Advisor: Prof. Zhihua Zhang;
- Accelerating random kaczmarz algorithm via clustering;
- Improving Block-RPCA algorithm performance on image inpainting.

PROFESSIONAL SERVICES

- **Associate Editor:** TVCG;
- **Conference Area Chair:** CVPR 2025; CoRL 2025; NeurIPS 2022, 2023, 2024, 2025;
- **Technical Program Committee:** SIGGRAPH Asia 2024;
- **International Program Committee:** Eurographics 2024, 2025;
- **Senior Program Committee Member:** AAAI 2023, 2024, 2025;
- **Workshop/Tutorial Organizers:** The 2nd Point Cloud Tutorial “All You Need To Know About 3D Point Cloud” [CVPR 2025]; AI for 3D Generation (AI3DG) [CVPR 2024]; AI for 3D Content Generation (AI3DCC) [ICCV 2023, ECCV 2024,]; Structural and Compositional Learning on 3D Data (StruCo3D) [ICCV 2021, CVPR 2023]; Visual Object-oriented Learning meets Interaction: Discovery, Representations, and Applications (VOLI) [ECCV 2022]; 3D Vision and Robotics (3DVR) [CVPR 2023]; The Rhobin Challenge – Reconstruction of human-object interaction [CVPR 2024, CVPR 2023];
- **Conference Reviewer:** ICML 2021, 2022, 2023; ICLR 2021, 2022, 2023, 2024; NeurIPS 2020, 2021, 2022, 2023; NeurIPS Datasets and Benchmarks Track 2021; CVPR 2020, 2021 (outstanding reviewer), 2022, 2023, 2024; ICCV 2019, 2021, 2023; ECCV 2020, 2022, 2024; CoRL 2022, 2024; RSS 2021, 2023; ICRA 2020, 2024; IROS 2021, 2024; Siggraph 2021, 2022, 2023, 2024; Siggraph Asia 2020, 2021, 2022, 2023; AAAI 2020, 2021, 2022; ACCV 2020; 3DV 2017, 2018, 2019, 2020, 2021, 2022; Pacific Graphics 2020; WACV 2020, 2021, 2022; MVA 2019;
- **Journal Reviewer:** RA-L; TPAMI; TVCG; TIP; TRO; TOG.

MEDIA COVERAGE

- **Computer Vision News (June 2022):** *PhD Dissertation Featuring;*
- **VentureBeat:** *Stanford researchers propose AI that figures out how to use real-world objects;*
- **IEEE Spectrum:** *Massive 3D dataset helps robots understand what things are;*

- **The Robot Report:** *Intel, OSU, Stanford, and UC San Diego work on reinforcement learning, PartNet could help household robots;*
- **Robotics Business Review:** *Researchers Launch 26K+ Object Dataset to Help Robots Learn Shapes;*
- **TechCrunch:** *Intel is doing the hard work necessary to make sure robots can operate your microwave;*
- **Intel AI Blog:** *Introducing PartNet: the first large-scale dataset with fine-grained, hierarchical, instance-level part annotations.*

HONORS AND AWARDS

- **School of Engineering Fellowship**, Stanford, 2016-2017
- **Meritorious Winner**, 2015 Mathematical Contest In Modeling, 2015 (top 10% of all participants)
- **National Scholarship**, 2015 (highest honor for undergraduates in China, awarded to top 1% students)
- **National Scholarship**, 2014 (highest honor for undergraduates in China, awarded to top 1% students)
- **KoGuan Scholarship**, 2013 (awarded to top 12 students in Zhiyuan College)
- **The First Prize**, National High School Mathematics Contest, China, 2011
- **The First Prize**, National Olympiad in Informatics in Provinces, China, 2010

TEACHING EXPERIENCE

- Guest Lecturer, Spring 2022, Machine Learning for 3D Data (CS 492A), KAIST.
- Teaching Assistant, Winter 2022, Neural Generative Models for 3D Geometry (CS 348n), Stanford University
- Guest Lecturer, Spring 2021, Geometric and Topological Data Analysis (CS 233), Stanford University
- Guest Lecturer, Winter 2021, Machine Learning Meets Geometry (CSE 291-I00), UCSD
- Teaching Assistant, Spring 2020, Geometric and Topological Data Analysis (CS 233), Stanford University
- Guest Lecturer, Spring 2018, Geometric and Topological Data Analysis (CS 233), Stanford University
- Teaching Assistant, Fall 2014, Introduction To Computer Science (CS 120), Shanghai Jiao Tong University

INVITED TALKS

- May 2022, *Learning Compositional and Actionable Visual Representations for 3D Shapes*, MIT Vision and Graphics Seminar.
- May 2022, *Compositional and Structural Learning on 3D Shapes*, Meta AI.
- Mar 2022, *Learning Compositional, Structural, and Actionable Visual Representations for 3D Shapes*, Google Research.
- Mar 2022, *Learning Actionable and Compositional Visual Representations for 3D Shapes*, Univ. of Washington.
- Feb 2022, *Learning Actionable and Compositional Visual Representations for 3D Shapes*, at Seattle Robotics Lab, NVIDIA.
- Dec 2021, *Self-supervised Affordance Learning for Manipulating Articulated 3D Objects*, at 4th Robot Learning Workshop, NeurIPS 2021.
- Dec 2021, *Learning 3D Shape Structure and Semantics*, at CompVision Seminar, UC Berkeley.
- Nov 2021, *Learning 3D Shape Structure and Semantics*, at Machine Learning and Friends Lunch, UMass Amherst.
- Aug 2021, *Learning 3D Shape Visual Actionable Information for Robotic Manipulation*, at SVL Reading Group, Stanford.
- June 2021, *Learning 3D Shape Structure and Semantics*, at CogAI Reading Group, Stanford.
- April 2021, *Learning 3D Shape Actionable Information from Simulated Interaction*, at SFU VCR (visual computing and robotics) seminar.
- March 2021, *Learning 3D Shape Actionable Information from Simulated Interaction*, at Autodesk Research.
- Feb 2021, *Where2Act: From Pixels to Actions for Articulated 3D Objects*, for the MatchLab at Imperial College London.

- Feb 2021, *Where2Act: From Pixels to Actions for Articulated 3D Objects*, for the PRIOR team at AI2.
- June 2020, *Part-level and Structural 3D Shape Understanding*, at Intelligent Graphics Laboratory (IGL).
- June 2020, *Part-level and Structural Understanding for 3D shape Perception, Synthesis and Editing*, GAMES: Graphics And Mixed Environment Seminar.
- April 2019, *PartNet: A Large-scale Benchmark for Fine-grained and Hierarchical Part-level 3D Object Understanding*, at Stanford GCafe Seminar.